

Summer of Science 2022



Operating Systems

Jaideep Kotani

200070035

Mentor: Oshin

June 2022

Contents

1	Introduction	3
2	Introduction to Operating Systems	3
2.1	What is an OS?	3
2.2	Background of OS	4
2.3	Design Goals of OS	4
3	Process Abstraction	5
3.1	Constitutes of a Process	5
3.2	Creating a Process	5
3.3	Process State Transitions	5
4	Process API	5
4.1	System Calls in Linux	7
4.1.1	Fork()	7
4.2	Exec()	8
5	Mechanism of Process Execution	8
6	Scheduling Policies	9
7	Inter Process Communication (IPC)	9
7.1	Shared Memory	9
7.2	Signals	9
7.3	Sockets	10
7.4	Pipes	10
8	Introduction to Virtual Memory	10
8.1	Allocating Memory	10
9	Address Translation	11
9.1	Role of OS in Translation	11
10	Paging	11
11	Memory Allocation Algorithms	12
11.1	Variable sized allocation:	12
11.2	Free List	13
11.3	Variable Size Allocation Strategies	13
12	Threads and Concurrency	13
12.1	Single Threaded Process	14
12.2	Multi threaded process	14
12.3	Threads	14
12.4	Scheduling Threads	14

13 Locks	14
13.1 Building Locks	14
13.2 Usage of Locks	15
14 Condition Variables	15
15 Semaphores	15
16 Concurrency Bugs	15
16.1 Deadlock Bugs	16
16.2 Non Deadlock Bugs	16
17 Files and Directories	16
17.1 The File Abstraction	16
17.2 Directory Tree	16
17.3 Operations	17
18 File System Implementation	17
18.1 File System	17
18.2 Inode Table	18
18.3 Inode Structure	18
18.4 File Allocation Table (FAT)	18
18.5 Directory Structure	18
18.6 Free Space Management	18
18.7 Virtual File System	19
18.8 Disk Buffer Cache	19

1 Introduction

This report is on topic **Operating Systems** which is my SoS project under MnP club, Cosmology and DarkMatter in the year 2022 and my mentor is Oshin. So in this half term of the course I allotted my time to learn about what OS is and following are the references I followed

- Lecture Series by Prof. Mythili Vutukuru
- Operating Systems: Three Easy Pieces by Remzi H. Arpaci-Dusseau
- Operating System Concepts by Silberschatz, Galvin and Gagne
- And ofcourse google and youtube as well

In this report, I would briefly summarize what I have learned and the theory covered by me in the past month. I briefly shared basic knowledge about Operating Systems.

The OS is used everywhere in places such as colleges, banks, companies, electronic devices etc. No device can operate without an operating system because it controls all the user's commands. For example Android OS, iOS and Symbian OS are used in mobile phones which are very lightweight operating systems and Linux is used in banking sector because it very secure OS.

The growth of the operating system is commendable as it was developed in 1950 to handle storage tape.

2 Introduction to Operating Systems

2.1 What is an OS?

An Operating System is an interface between the user and computer hardware. An OS is a software which handles all tasks like

- Memory Management
- Processor Management
- Device Management
- Security
- Reliability
- File Management
- Job accounting

2.2 Background of OS

When a program is run a compiler translates the program into an executable which contains instructions that CPU processes and all data of program. The instructions that run on CPU are implemented by instruction set architecture (ISA). The executable generally fetches instruction pointed by PC from memory and loads needed data by instructions into registers. then it executes the given instructions by decoding them and stores the final results obtained in the memory.

What OS does basically here is it handles all the program memory, CPU and external devices.

Managing CPU

- creating and managing processes i.e process abstraction
- Virtualization of CPU
- time sharing between processes and the coordination between them

Managing Memory

- Basically it is managing of memory of process i.e code, data and stack etc
- Assigning virtual addresses
- Translation of virtual addresses to physical addresses

Managing Devices

- These devices include disk, network card, graphics card and other external devices
- Issuing instructions to devices and responding to interrupt events like keys of keyboard

2.3 Design Goals of OS

In view of User:

The OS should be easy to use, fast, reliable and safe. The OS should be designed in view of a general user whose user base is more. This difficult to achieve as there's no specific method to progress forward.

In view of System Design:

The OS should be easy to design, implement and maintain which are goals of the people who create and maintain the OS. These goals have many subset of goals like having ability to handle multiple users and devices parallelly, providing security and privacy, ability to weather future software and hardware changes, making OS portable and backward compatibility.

3 Process Abstraction

Abstraction provide an interface to application programmers that separates **mechanism** (implementation of interface) from **policy** (what interface commits to accomplish).

In process abstraction, details of the threads of execution are not visible to the consumer of the process. An example of process abstraction is the concurrency scheduler in a database system. A database system can handle many concurrent queries.

3.1 Constitutes of a Process

A unique Process Identifier(PID), memory image that comprises of static and dynamic memories, CPU context which includes various registers and file descriptors i.e pointers to open files and devices.

3.2 Creating a Process

While creating a process, OS assigns a PID to it and inserts a new entry in primary process table. Then required memory space for all elements of processes are allocated space in PCB(Process Control Block). After this various values in PCB are initialized. Then OS will link these processes to queue and now it competes for CPU. The OS will create data structures like log files to keep track of the activity of the processes.

States of Process

- Running - which is currently running on CPU
- Ready - waiting to be scheduled
- Blocked - suspended and not ready to run
- New - being created and yet to run
- Dead - terminated

3.3 Process State Transitions

The states that a Process enters in working from start till end are known as Process states. The minimum number of states is five.

4 Process API

API is Application Programming Interface which is a collection of communication protocols and subroutines used by various programs to communicate between them. Briefly defined as functions available to write user programs.

API provided by OS is a set of system calls which is a functional call into OS that runs at higher priority level of CPU. These are sensitive operations for




Figure 1: General State Diagram of Transitions

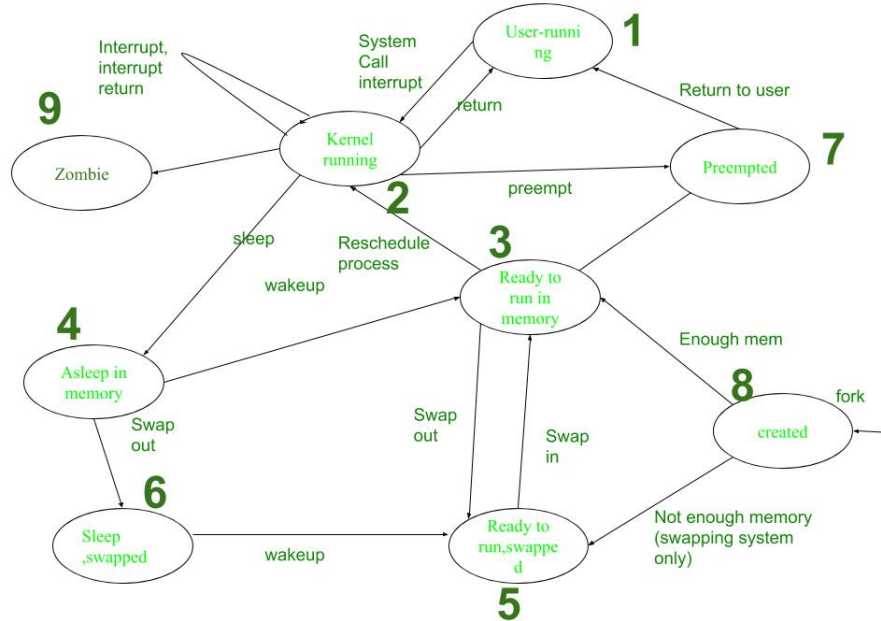


Figure 2: Transitions in a UNIX Process

example accessing hardware etc. Some system call may even cause process to be blocked and rescheduled.

4.1 System Calls in Linux

This a procedure that provides the interface between a process and the operating system. It is the way by which a computer program requests a service from the kernel of the operating system. Different operating systems execute different system calls.

In Linux, making a system call involves transferring control from unprivileged user mode to privileged kernel mode; the details of this transfer vary from architecture to architecture. The libraries take care of collecting the system-call arguments and, if necessary, arranging those arguments in the special form necessary to make the system call.

Process Control:

This system calls perform the task of process creation, process termination, etc. The Linux System calls under this are `fork()` , `exit()` , `exec()`.

4.1.1 Fork()

- A new process is created by making a copy of parent's memory image
- The new process is added to OS process list and scheduled
- Parent and Child start execution just after fork and modify the memory data independently

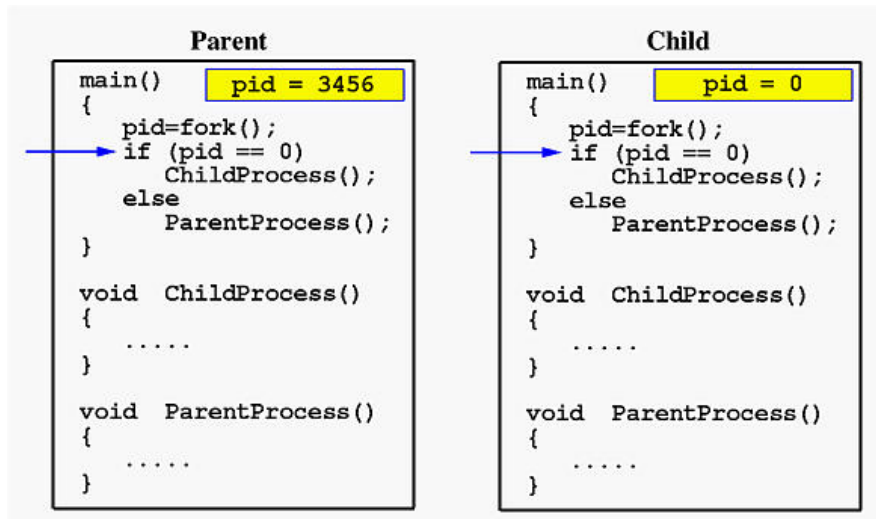


Figure 3: The `fork()` system call

4.2 Exec()

- After fork, parent and child are running same code
- A process can run `exec()` to load another executable to its memory image
- Variants of `exec()` are used to pass commandline arguments to new executable

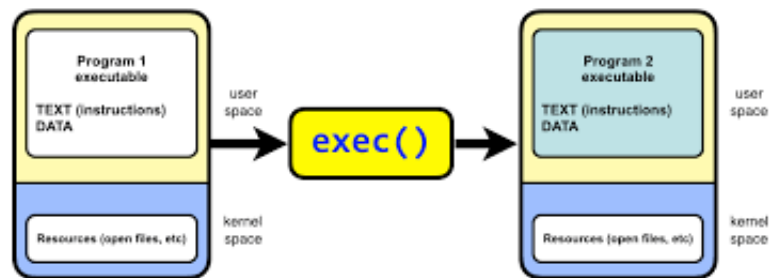


Figure 4: The `exec()` system call

5 Mechanism of Process Execution

Process Execution:

- OS allocates memory and creates memory image
- Points CPU PC to current instruction
- After setup, OS is out of way and process executes directly on CPU

OS Scheduler:

- The scheduler has 2 parts respectively named Policy and Mechanism
- Cooperative schedulers are polite and switch only if process is blocked
- non-cooperative schedulers can switch even when process is ready to continue where CPU generates periodic timer interrupts
- After interrupt the OS checks if current process has run long enough or not

6 Scheduling Policies

CPU Scheduling is a process that allows one process to use the CPU while another process is delayed (in standby) due to unavailability of any resources such as I / O etc, thus making full use of the CPU. The purpose of CPU Scheduling is to make the system more efficient, faster, and fairer. **What are we optimizing?:**

- maximize the utilization
- minimize average time i.e time from process arrival to first scheduling
- minimize average response time
- fairness that all processes have same priority
- minimize overhead i.e run process long enough to reduce cost of context switch

Few policies include First-In-First-Out (FIFO), Shortest Job First (SJF), Shortest Time-to-Completion First (STCF), Round Robin (RR) etc. Each policy has its merits and demerits.

Example: Linux uses a Multi Level Feedback Queue (MLFQ)

7 Inter Process Communication (IPC)

This is a mechanism provided by the operating system that allows processes to communicate with each other. This communication could involve a process letting another process know that some event has occurred or the transferring of data from one process to another.

7.1 Shared Memory

Processes can both access same region of memory via "shmget()" system call. By providing same key, 2 processes can get same segment of memory. They can read or write to memory to communicate.

7.2 Signals

A certain set of signals are supported by OS. Signals can be sent to a process by OS. There's a signal handler where every process has a default code to execute for each signal. Some signal handlers can be overridden to do other things also.

7.3 Sockets

Sockets can be used for two processes on same machine or different machines to communicate. TCP or UDP sockets are present across machines. Communication with sockets is done when processes open sockets and connect them to each other. Messages written into one socket can be read from another. The OS transfers data across socket buffers

7.4 Pipes

Pipe system call returns two file descriptions which are read and write handles. A pipe is a half-duplex communication. The data written in one file descriptor can be read through another.

There are two types of pipes **Regular Pipes** and **Named Pipes**. In regular pipes, parent and child share fd after fork where parent and child use respective ends of the pipe. In named pipes, two endpoints of a pipe can be in different processes. Pipe data buffered in OS buffers between write and read.

8 Introduction to Virtual Memory

Virtual Memory is a storage allocation scheme in which secondary memory can be addressed as though it were part of the main memory. The addresses a program may use to reference memory are distinguished from the addresses the memory system uses to identify physical storage sites, and program-generated addresses are translated automatically to the corresponding machine addresses.

The size of virtual storage is limited by the addressing scheme of the computer system and the amount of secondary memory is available not by the actual number of the main storage locations. Now multiple active processes timeshare CPU. We are using virtual memory because the real view of memory is very messed up. **Goals of Virtualization**

- Transparency: user programs should not be aware of the messy details
- Efficiency: minimize overhead and wastage in terms of memory space and access time
- Isolation and protection: a user process should not be able to access anything outside its address space

8.1 Allocating Memory

OS allocates a set of pages to the memory image of the process. Within the image static variables are allocated in the executable. Local variables of a function are on stack and Dynamic allocation is done with "malloc" on the heap

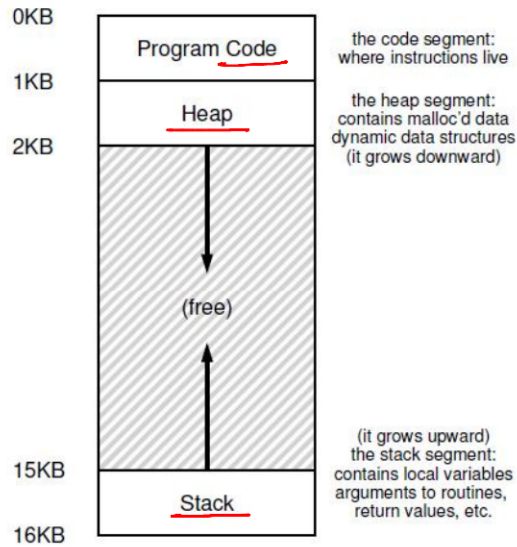


Figure 5: Address Space

9 Address Translation

Virtual address space is setup by OS during process creation. A simplifies OS places entire memory image into one chunk. CPU provides the mode of execution in translation. ISA has instructions that set translation information.

Hardware uses this information to perform translation on every memory access. This hardware generates faults and traps to OS when access is illegal.

9.1 Role of OS in Translation

- OS maintains free list of memory
- Allocates space to process during creation and cleans up when done
- Maintains information of where space is allocated to each process in PCB
- Sets address translation information in hardware and updates this information upon context switch

10 Paging

Paging is a memory management scheme that eliminates the need for contiguous allocation of physical memory. This scheme permits the physical address space of a process to be non-contiguous.

- Logical Address or Virtual Address : An address generated by the CPU
- Logical Address Space or Virtual Address Space: The set of all logical addresses generated by a program
- Physical Address : An address actually available on memory unit
- Physical Address Space : The set of all physical addresses corresponding to the logical addresses

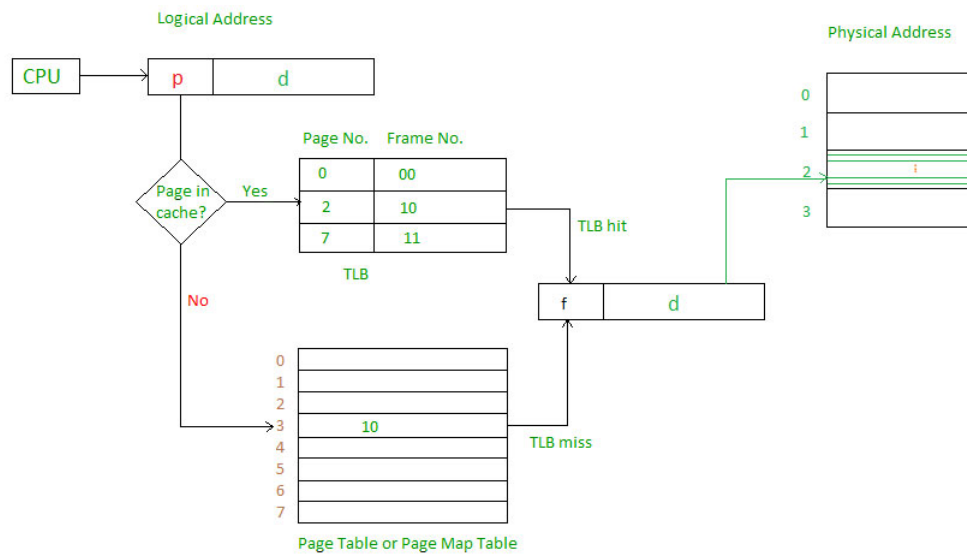


Figure 6: Paging

11 Memory Allocation Algorithms

In the operating system, the following are four common memory management techniques. Single contiguous allocation, Partitioned allocation, Paged memory management, Segmented memory management. Most of the operating systems (for example Windows and Linux) use Segmentation with Paging. A process is divided into segments and individual segments have pages.

11.1 Variable sized allocation:

Considering a simple implementation of "malloc", every allocated chunk has a header with info like size of chunk.

11.2 Free List

Free space is managed as a list. Pointer to the next free chunk is embedded within the free chunk. The library tracks the head of the list – Allocations happen from the head.

11.3 Variable Size Allocation Strategies

- First fit: allocate first free chunk that is sufficient
- Best fit: allocate free chunk that is closest in size
- Worst fit: allocate free chunk that is farthest in size

12 Threads and Concurrency

A thread is a path of execution within a process. A process can contain multiple threads. A thread is also known as lightweight process. The idea is to achieve parallelism by dividing a process into multiple threads.

For example, in a browser, multiple tabs can be different threads. MS Word uses multiple threads: one thread to format the text, another thread to process inputs, etc.

Concurrency is the execution of the multiple instruction sequences at the same time. It happens in the operating system when there are several process threads running in parallel. The running process threads always communicate with each other through shared memory or message passing.

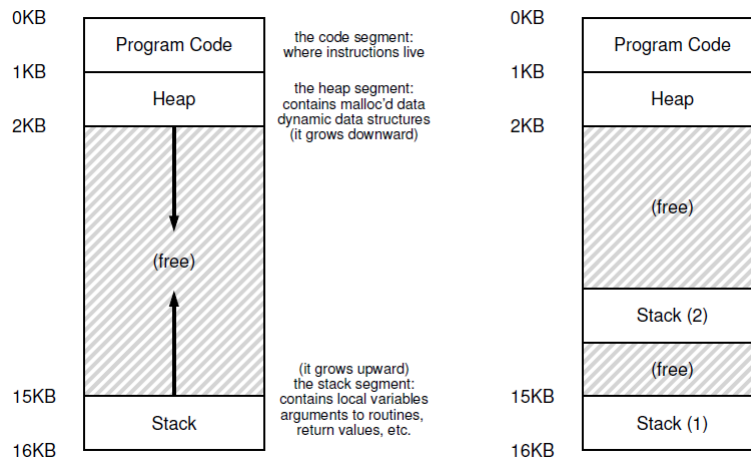


Figure 7: Single-Threaded And Multi-Threaded Address Spaces

12.1 Single Threaded Process

Single threaded processes contain the execution of instructions in a single sequence. In other words, one command is processed at a time. The opposite of single threaded processes are multithreaded processes. These processes allow the execution of multiple parts of a program at the same time. These are lightweight processes available within the process.

12.2 Multi threaded process

Multithreading is the ability of a program or an operating system to enable more than one user at a time without requiring multiple copies of the program running on the computer. Multithreading can also handle multiple requests from the same user.

12.3 Threads

A single process can effectively utilize multiple CPU cores which is known as Parallelism. Concurrency means running multiple threads/processes at the same time, even on single CPU core, by interleaving their executions. Parallelism means running multiple threads/processes in parallel over different CPU cores.

12.4 Scheduling Threads

Scheduling of threads involves two boundary scheduling:

- Scheduling of user level threads (ULT) to kernel level threads (KLT) via lightweight process (LWP) by the application developer.
- Scheduling of user level threads (ULT) to kernel level threads (KLT) via lightweight process (LWP) by the application developer.

13 Locks

We use a special lock variable to protect shared variables. All threads accessing a critical section share a lock. One thread succeeds in locking – owner of lock. Other threads that try to lock cannot proceed further until lock is released by the owner. Pthreads library in Linux provides such locks.

13.1 Building Locks

Lock Implementation:

- Mutual Exclusion and implementing needs support from OS and hardware
- All threads should eventually get the lock, and no thread should starve
- Acquiring, releasing, and waiting for lock should not consume too many resources i.e low overhead

13.2 Usage of Locks

- A lock should be acquired before accessing any variable or data structure that is shared between multiple threads of a process
- All shared kernel data structures must also be accessed only after locking
- Fine-grained allows more parallelism
- Multiple fine-grained locks may be harder to manage

14 Condition Variables

A condition variable (CV) is a queue that a thread can put itself into when waiting on some condition. Another thread that makes the condition true can signal the CV to wake up a waiting thread. Pthreads provides CV for user programs. Signal wakes up one thread, signal broadcast wakes up all waiting threads.

15 Semaphores

Synchronization primitive like condition variables. Semaphore is a variable with an underlying counter. There are two functions on a semaphore variable:

- Up/post increments the counter
- Down/wait decrements the counter and blocks the calling thread if the resulting value is negative
- A semaphore with init value 1 acts as a simple lock

16 Concurrency Bugs

A concurrency bug is an (undesired) outcome that arises if two programs are run at the same time that does not show up if the two programs are run sequentially, one after the other. Common bugs in concurrent programs:

- Bugs are non-deterministic and occur based on execution order of threads
 - very hard to debug
- Deadlocks: threads cannot execute any further and wait for each other
- Non-deadlock bugs: non deadlock but incorrect results when threads execute

16.1 Deadlock Bugs

In an operating system, a deadlock occurs when a process or thread enters a waiting state because a requested system resource is held by another waiting process, which in turn is waiting for another resource held by another waiting process.




Figure 8: The Deadlock Dependency Graph

16.2 Non Deadlock Bugs

- Atomicity bugs: atomicity assumptions made by programmer are violated during execution
- Order-violation bugs: desired order of memory accesses is flipped during concurrent execution of concurrent threads

17 Files and Directories

17.1 The File Abstraction

- File is a linear array of bytes that is stored persistently. It is identified with file name and a OS-level identifier. Inode number unique within the file system.
- Directory contains other subdirectories and files along with their inode numbers. These are stored like a file whose contents are filename-to-inode mappings

17.2 Directory Tree

By placing directories within other directories, users are able to build an arbitrary directory tree (or directory hierarchy), under which all files and directories are stored.




Figure 9: An Example Directory Tree

17.3 Operations

Operations on Files:

- Creating a file and Opening a file
- Reading/writing files
- Close system call
- Seek and Sync

Operations on Directories are similar to as for files, and they can be accessed by create, open, read, close.

18 File System Implementation

18.1 File System

- An organization of files and directories on disk
- OS has one or more file systems
- Data structures to organize data and metadata on disk
- Implementation of system calls like open, read, write using the data structures
- Disks expose a set of blocks. System Calls translated into reads and writes on blocks

18.2 Inode Table




Figure 10: Inode Table

Usually Inodes are stored in array. Inode number of file is index into this array. Inode stores file metadata like permissions, access time, etc and pointers of file data i.e disk block numbers.

18.3 Inode Structure

File data not stored contiguously on disk, need to track multiple block numbers of a file. Inode tracking disk block numbers:

- Direct pointers: numbers of first few blocks are stored in inode itself (suffices for small files)
- Indirect block: for larger files, inode stores number of indirect block, which has block numbers of file data.
- In the same way double and triple indirect blocks are tracked.

18.4 File Allocation Table (FAT)

This is an alternate way to track file blocks. FAT stores next block pointer for each block. FAT has 1 entry per disk block. Entry has number of next file block or null. Pointer to first block is stored in inode.

18.5 Directory Structure

- Directory stores records mapping filename to inode number
- Linked list of records, or more complex structures like hash tables, binary search trees etc.
- Directory is a special type of file and has inode and data blocks

18.6 Free Space Management

We track free blocks using Bitmaps, Free list and other complex structures.

- Bitmaps: For inodes and data blocks, we store 1 bit per block to indicate if free or not

- Free List: Super block stores pointer to first free block, a free block stores address of next block on list

18.7 Virtual File System

A virtual file system (VFS) is programming that forms an interface between an operating system's kernel and a more concrete file system. The VFS serves as an abstraction layer that gives applications access to different types of file systems and local and network storage devices

- File systems differ in implementations of data structures (e.g., organization of file records in directory)
- Linux supports virtual file system (VFS) abstraction
- VFS looks at a file system as objects (files, directories, inodes, superblock) and operations on these objects (e.g., lookup filename in directory)
- System call logic is written on VFS objects
- To develop a new file system, simply implement functions on VFS objects and provide pointers to these functions to kernel
- Syscall implementation does not have to change with file system implementation details

18.8 Disk Buffer Cache

Disk buffer is the embedded memory in a hard disk drive (HDD) acting as a buffer between the rest of the computer and the physical hard disk platter that is used for storage. It is physically distinct from and is used differently from the page cache typically kept by the operating system in the computer's main memory. It is controlled by the microcontroller in the hard disk drive, and the page cache is controlled by the computer to which that disk is attached.

- Results of recently fetched disk blocks are cached
- File system issues block read/write requests to block numbers via buffer cache
 - If block in cache, served from cache, no disk I/O
 - If cache miss, block fetched to cache and returned to file system
- Writes are applied to cache block first
 - Synchronous/write-through cache writes to disk immediately
 - Asynchronous/write-back cache stores dirty block in memory and writes back after a delay

- Unified page cache in OS where free pages allocated to both processes and cache from common pool

Benefits of disk buffer cache are that there is improved performance due to reduced disk I/O and there is single copy of block in memory which gives consistency across processes